
Generalised Off-line Training of Spiking Deep Neural Networks

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Spiking neural networks (SNNs) can be trained by first training an equivalent ANN
2 and then transferring the trained weights to the SNN, but there are two significant
3 problems to be solved. First, an accurate activation function is needed to model
4 the neural dynamics of spiking neurons, and our previously proposed activation
5 function, Noisy Softplus, has shown to be a good match to the response activity
6 of Leaky Integrate-and-Fire (LIF) neurons. The second problem is mapping the
7 abstract numerical values of the ANN to concrete physical units in the SNN, such
8 as current and firing rate. In this paper, we introduce the parametric activation
9 function (PAF) to tackle the second problem. With these problems solved, SNNs
10 can be trained exactly the same way as ANNs, and the trained weights can be
11 used directly in the spiking version of the same network without any conversion.
12 More importantly, the PAF can be generalised to activation functions other than
13 Noisy Softplus, such as the Rectified Linear Unit (ReLU). Therefore, SNN training
14 can be simplified as: estimate the parameters for PAF according to biological
15 configurations of LIF neurons, and use parametric ReLU as the activation function
16 when training. In addition, we propose a fine tuning method as an option which
17 helps the trained network to match the SNN more closely. Based on this generalised
18 training method, we achieve the best SNN accuracy on the MNIST task using LIF
19 neurons, 98.85%, on a 6-layer spiking convolutional neural network (ConvNet).

20 1 Introduction

21 Advances in computing power and deep learning have benefited computers with a rapid growing
22 performance in cognitive tasks, such as recognising objects [1] and playing GO [2]). These tasks
23 were once dominated by human intelligence and solved by biological neurons in the brain. However,
24 humans and many other animals still win against computers in practical tasks and outperform in terms
25 of size and energy cost over several orders of magnitude. For instance, AlphaGO [2] consumed 1 MW
26 of power on its 1920 CPUs and 280 GPUs when playing the game with one of the best human players
27 whose brain only consumed about 20 W. Although we are still far from understanding the brain
28 thoroughly, it is believed that the performance gap between computation in the biological nervous
29 system and in a computer lies in the fundamental computing units and the way they act. Computers
30 employ Boolean logic and deterministic digital operations based usually on synchronous clocks while
31 nervous systems employ parallel-distributed, event-driven, stochastic unreliable components [3].
32 The impressive disparities in cognitive capabilities and energy consumption drives the research into
33 biologically-plausible spiking neurons.

34 A regular artificial neuron (Figure 1(a)) composes a weighted summation of input data, $\sum x_i w_i$,
35 and an activation function, f , applying on the sum. Usually, a bias is included in the weighted
36 summation which can be seen as an extra input $x_b = 1$ with its weight set to b . However, in this




Figure 1: Comparisons of processing mechanisms of an artificial and a spiking neuron. (a) An artificial neuron takes numerical values of vector \mathbf{x} as input, works as a weighted summation followed by an activation function f . (b) Spike trains flow into a spiking neuron as input stimuli, trigger linearly summed-up PSPs through synapses with different synaptic efficacy \mathbf{w} , and post-synaptic neuron generates output spikes when membrane potential reaches some threshold.

paper we exclude biases for both artificial and spiking neurons to simplify neural models and to reduce parameters. Meanwhile the inputs of a spiking neuron (Figure 1(b)) are pre-synaptic spike trains, which create Post-synaptic Potentials (PSPs) and trigger spikes as outcomes when the neuron's membrane potential reaches some threshold. Neural dynamics of the membrane potentials, PSPs, and spike trains are all time dependent, while neurons of ANNs, e.g. sigmoid units, only cope with abstract numerical values representing spiking rate, without timing information. These fundamental differences on input/output representation and neural computation raise the research problem of how to operate and train biologically-plausible SNNs to be competent in cognitive tasks just as ANNs.

An intuitive idea is to train SNNs on an equivalent ANN and then transfer the trained weights to the SNN. Jug et al[4] first proposed the use of the Siegert formula as the activation function in training Deep Belief Networks, which maps incoming currents driven by Poisson spike trains to the response firing rate of a LIF neuron. The activation function is in actual physical units, thus the trained weights can be transferred directly into SNNs. However, most importantly, the numerical analysis on an LIF response function using the Siegert formula is far from accurate. Moreover, the high complexity of the Siegert function causes much longer training time and more energy, let alone the high-cost computation on its derivative to obtain the error gradient.

Hence researchers turned to abstract activation functions to model the response activity of spiking neurons to simplify the computation. This identifies the first problem of accurately model the neural dynamics of LIF neurons with abstract activation functions. Recent work [13] proposed the Soft LIF response function, which is equivalent to Softplus activation, but the algorithm models the firing activity with even worse accuracy by getting rid of the current noise. Noisy softplus [15], on the other hand, includes noise as the second factor, takes account of biological parameters, and is proved to be a close match to the response activity of LIF neurons.

Then the following problem appeared to map the abstract activation functions to actual physical units: current in nA and firing rates in Hz . Instead of solving the problem, the alternative way of converting ANN-trained weights to be fitted in SNNs [8, 9] was successfully applied on less biologically-realistic and simplified integrate-and-fire (IF) neurons. Normalising the trained weights for use on simplified IF neurons was relatively straightforward, so this method keeps the state-of-the-art performance. Therefore, the paper aims to address the second problem of mapping activation functions to actual response firing activity of LIF neurons, thus to complete the generalised SNN training mechanism.

2 ANN-Trained SNNs

2.1 Extended Noisy Softplus

Suppose λ_{out} represents the firing rate of a LIF neuron driven by a noisy current x , we extended Noisy Softplus[15] with a scale factor, S , then

$$\begin{aligned} \lambda_{out} &\simeq f_{ns}(x, \sigma) \times S \\ &= k\sigma \log[1 + \exp(\frac{x}{k\sigma})] \times S. \end{aligned} \quad (1)$$

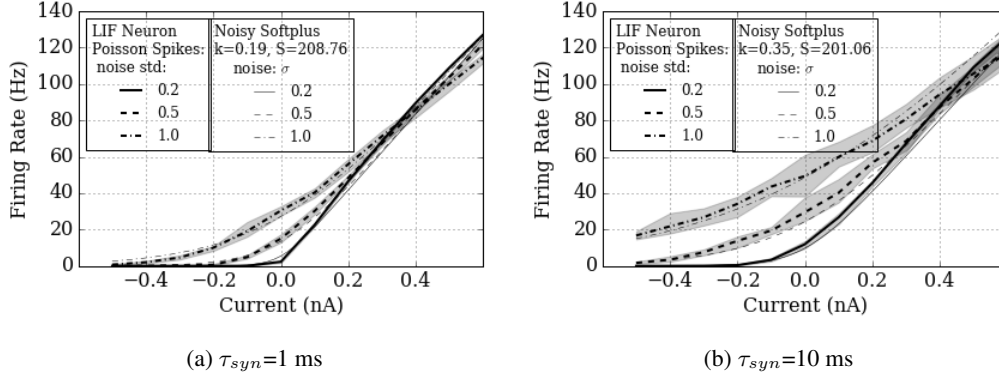


Figure 2: Noisy Softplus fits to the response firing rates of LIF neurons.

Noisy Softplus fits well to the practical response firing rate of the LIF neuron with suitable calibration of k and S , see Figure 2. The parameter pair of (k, S) is curve-fitted with the triple data points of (λ, x, σ) . The fitted parameter was set to $(k, S) = (0.19, 208.76)$ for the practical response firing rate driven by synaptic noisy current with $\tau_{syn} = 1$ ms and was set to $(k, S) = (0.35, 201.06)$ when $\tau_{syn} = 10$ ms. The calibration currently is conducted by linear least squares regression; numerical analysis is considered however for future work to express the factors with biological parameters of a LIF neuron. From now on, we can model the response firing activity of a LIF neuron with a unified activation function, extended Noisy Softplus.

In order to illustrate how we can use the extended Noisy Softplus to train layered up SNNs, we will demonstrate the mapping between the physical activity and numerical ANN calculations in the following subsections.

2.2 Equivalent Input and Output

Neurons in ANNs take inputs from their previous layer, and feed the weighted sum of their input, $net_j = \sum_i w_{ij}x_i$, to the activation function. The transformed signal then forms the output of an artificial neuron, which can be denoted as $y_j = f(net_j)$. However, Noisy Softplus takes physical quantities of current, and firing rate as input/output, thus an extra step is still needed to map the firing rate to numerical values in ANNs. According to Equation [15], the mean of the current feeding into a spiking neuron is equivalent to net of artificial neurons, where

$$m_{I_j} = \sum_i w_{ij}(\lambda_i \tau_{syn}), \text{ then} \quad (2)$$

$$net_j = \sum_i w_{ij}x_i, \text{ and } x_i = \lambda_i \tau_{syn}.$$

The noise level of Noisy Softplus, σ^2 , is the variance of the current, which also can be seen as a weighted sum of the same input x but with different weights:

$$s_{I_j}^2 = \sum_i \left(\frac{1}{2} w_{ij}^2\right) (\lambda_i \tau_{syn}), \text{ then} \quad (3)$$

$$\sigma_j^2 = \sum_i \left(\frac{1}{2} w_{ij}^2\right) x_i.$$

Noisy Softplus transforms the noisy current with parameters of (net_j, σ_j) to the equivalent ANN output y_j , where it can be scaled up by the factor S to the firing rate of SNNs. Note that the calculation of noise level is not necessary for activation functions other than Noisy Softplus, for example, it can be set to a constant for Softplus or 0 for ReLU. We name the neuron model ‘artificial spiking neurons’ which takes firing rates of spike trains as input and output. The entire artificial spiking neuron model is then generalised to any ReLU/Softplus-like activation functions, See Figure 3.




Figure 3: Artificial spiking neuron takes scaled firing rates as input, then transforms weighted sum in some activation unit to its output which can be scaled-up to the firing rate of an output spike train.

2.3 Layered-up Network

Referred to Figure 3, if we move the left end process of $\times \tau_{syn}$ to the right end after λ_j , Figure 3 forms the same neuron model and structure as multilayer perceptron: neurons take x as input and outputs y , and this conversion is illustrated in Figure 4. The process within such an artificial neuron is divided into weighted summation and activation, which also applies to SNN modelling by combining the scaling factor S and the synaptic time constant τ_{syn} to activation functions. Thus the combined activation function for modelling SNNs should be:

$$y = f(x) \times S \times \tau_{syn} , \quad (4)$$

and its derivative function which is used when back propagates is:

$$\frac{\partial y}{\partial x} = f'(x) \times S \times \tau_{syn} . \quad (5)$$




Figure 4: Transforming artificial spiking neurons to artificial neurons for SNN modelling. The combined activation links the firing activity of a spiking neuron to the numerical value of ANNs.

Thus, using this method of ANN-trained SNNs, the activation functions are of lower complexity than the Siegert formula, and their corresponding derivative functions can be directly used for back propagation. Furthermore, the method enables ReLU-like activation functions for SNN training, thus improving the recognition accuracy while keeping a relative lower firing rate compared to sigmoid neurons. Most significantly, the ANN-trained weights are ready for use in SNNs without any transformation, and the output firing rate of a spiking neuron can be estimated in the ANN simulation.

2.4 Fine Tuning

There are two aspects to the fine tuning which makes the ANN closer to SNNs: Firstly, using Noisy Softplus activation functions in a whole trained network operates every single neuron running in a similar noise level as in SNNs, thus the weights trained by other activation functions will be tuned to fit closer to SNNs. Secondly, the output firing rate of any LIF neuron is greater than zero as long as noise exists in their synaptic input. Thus adding up a small offset on the labels directs the model to approximate to practical SNNs.

The labels of data are always converted to binary values for ANN training. This enlarges the disparities between the correct recognition label and the rest to train the network for better classification capability. Consequently, we can train the network with any activation function and then fine-tune it with Noisy Softplus to take account of both accuracy and practical network activities of SNNs. However, we add a small number, for example 0.01, to all the binary values of the data labels. Doing so helps the training to loosen the strict objective function to predict exact labels with binary values. Instead, it allows a small offset to the objective. An alternative method is to use Softmax function at the top layer, which aims to map real vectors to the range of (0, 1) that add up to 1. However, without a limit on the input of Softmax, it will be easy to reach or even exceed the highest firing rate of a spiking neuron. The result of fine tuning on a Convnet will be demonstrated in subsection 3.2.

3 Results

A convolutional network model was trained on MNIST, a popular database in neuromorphic vision, using the ANN-trained SNN method stated above. The architecture contains 28×28 input units, followed by two convolutional layers 6c5-2s-12c5-2s, and 10 output neurons fully connected to the last pooling layer to represent the classified digit.

The training only employed Noisy Softplus units that all the convolution, average sampling, and the fully-connected neurons use Noisy Softplus function with no bias. The parameters of the activation function were calibrated as, ($k = 0.30, S = 201$), for LIF neurons ($C_m = 0.25\text{nF}, \tau_m = 20.0\text{ms}, \tau_{refrac} = 1.0\text{ms}, v_{reset} = -65.0\text{mV}, v_{rest} = -65.0\text{mV}, v_{thresh} = -50.0\text{mV}, i_{offset} = 0.1\text{nA}, \tau_{syn} = 5\text{ms}$). The input images were scaled by 100 Hz to present the firing rates of input spikes. The weights were updated using a decaying learning rate, 50 images per batch and 20 epochs. The ANN-trained weights were then directly applied in the corresponding convolutional SNN without any conversion for recognition tasks.

3.1 Neural Activity

To validate how well the Noisy Softplus activation fits to the response firing rate of LIF neurons in a real application, we simulated the model on NEST using the Poisson MNIST dataset [16] and the neurons of a convolutional map were observed.

A small test of ten MNIST digits presented in Poisson spike trains for 1 s each. A trained 5×5 kernel was convolved with these input digits, and the convolved output of the feature map, the output firing rate was recorded during a real-time SNN simulation on NEST, and compared to the modelled activations of Equation 4 in ANNs.

The input x of the network was calculated as Equation 2: $x_i = \lambda_i \tau_{syn}$, and so as the weighted sum of the synaptic current (see Equation 2), net_j and its variance (see Equation 3), σ_j^2 . With three combined activation functions as Equation 4:

$$\begin{aligned} (1) \text{ Noisy Softplus: } y_j &= k\sigma_j \log[1 + \exp(\frac{net_j}{k\sigma_j})] \times S \times \tau_{syn} , \\ (2) \text{ ReLU: } y_j &= \max(0, net_j) \times S \times \tau_{syn} , \\ (3) \text{ Softplus: } y_j &= k\sigma \log[1 + \exp(\frac{net_j}{k\sigma})] \times S \times \tau_{syn} , \quad \sigma = 0.45, \end{aligned} \tag{6}$$

we compare the output to the recorded SNN simulations. ReLU assumes a non-noise current, and Softplus takes a static noise level thus σ_j is not used for either of them, meanwhile Noisy Softplus adapts to noise automatically with σ_j . The experiment took the sequence of 10 digits to the same kernel and the estimated spike counts using Noisy Softplus fit to the real recorded firing rate much more accurately than ReLU and Softplus, see 5. The Euclidean distance, $\sqrt{\sum_j (y_j/\tau_{syn} - \lambda_j)^2}$, between the spike counts and the predicted firing rates by Noisy Softplus, ReLU and Softplus was 184.57, 361.64 and 1102.76 respectively. We manually selected a static noise level of 0.45 for Softplus, whose estimated firing rates located roughly on the top slope of the real response activity. This resulted in longer Euclidean distance than using ReLU, since most of the input noisy currents were of relatively low noise level in this experiment. Hence, the firing rate driven by lower noise level is closer to ReLU curve than Softplus.

The SNN successfully classified the digits where the correct label neuron fired the most. We trained the network with binary labels on the output layer, thus the expected firing rate of correct classification was $1/\tau_{syn} = 200$ Hz according to Equation 4. The firing rates of the recognition test fell to the valid range around 0 to 200 Hz. This shows another advantage of the proposed ANN-trained method that we can constrain the expected firing rate of the top layer, thus preventing SNN from exceeding its maximum firing rate, for example 1000 Hz when time resolution of SNN simulation set to 1 ms.

3.2 Recognition Performance

Here we focus on the recognition performance of the proposed ANN-trained SNN method. Before looking into the recognition results, it is significant to see the learning capability of the proposed activation function, Noisy Softplus. We compared the training using ReLU, Softplus, and Noisy Softplus by their loss during training averaged over 3 trials, see Figure 6. ReLU learned fastest with the lowest loss, thanks to its steepest derivative. In comparison, Softplus accumulated spontaneous firing rates layer by layer and its derivative may experience vanishing gradients during back propagation, which result in a more difficult training. Noisy Softplus performance lay between these two in terms of loss and learning speed. However, the loss stabilised fastest, which means a possible shorter training time.

The recognition test took the whole testing dataset of MNIST which contains 10,000 images. At first, all trained models were tested on the same artificial neurons as used for training in ANNs, and these experiments were called ‘DNN’ test since the network had a deep structure (5 layers). Subsequently, the trained weights were directly applied to SNN without any transformation, and these ‘SNN’ experiments tested their recognition performance on the NEST simulator. The LIF neurons had the same parameters as in training. The input images were converted to Poisson spike trains and presented for 1 s each. The output neuron which fired the most indicated the classification of an input image. Moreover, a ‘Fine tuning’ test took the trained model for fine tuning, and the tuned weights were tested on the same SNN environment. The tuning only ran for one epoch, 5% cost of the ANN training (20 epochs), using Noisy Softplus neurons with labels shifted for +0.01.

The classification errors for the tests are investigated in Table 1 and the averaged classification accuracy is shown in Figure 7. From DNN to SNN, the classification accuracy declines by 0.80%, 0.79% and 3.12% on average for Noisy softplus, ReLU and Softplus. The accuracy loss was caused by the mismatch between the activations and the practical response firing rates, see example in Figure 5, and the strict binary labels for Noisy Softplus and Softplus activations. Fortunately, the problem is alleviated by fine tuning which increased the classification accuracy by 0.38%, 0.19% and 2.06%, and resulted in the total loss of 0.43%, 0.61%, and 1.06% respectively. The improvement of ReLU is not as great as the others, because there is no problem of strict labels during training. Softplus benefits the most from fine tuning, since not only the huge mismatch of response firing rate is greatly corrected, but also the offset on the labels helps the network to fit SNNs.

The most efficient training in terms of both classification accuracy and algorithm complexity, takes ReLU for ANN training and Noisy Softplus for fine tuning. Softplus does not exhibit better classification capability and more importantly the manual selected static noise level hugely influences the mismatch between the predicted firing rates and the real data. Although Noisy Softplus shows the least classification drop from ANNs to SNNs, the training performance is still worse than ReLU.




Figure 5: Noisy Softplus fits to the neural response firing rate in an SNN simulation. The recorded firing rate of the same kernel convolved with 10 images in SNN simulation, comparing to the prediction of activations of Softplus, ReLU, and Noisy Softplus.




Figure 6: Comparisons of Loss during training using Noisy Softplus, ReLU and Softplus activation functions. Bold lines show the average of three training trials, and the grey colour illustrates the range between the minimum and the maximum values of the trials.


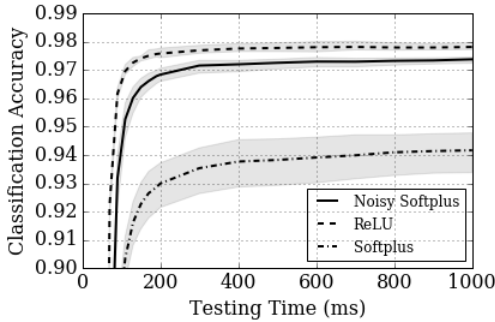


Figure 7: Classification accuracy compared among trained weights of Noisy Softplus, ReLU, Softplus on DNN, SNN and fine-tuned SNN.

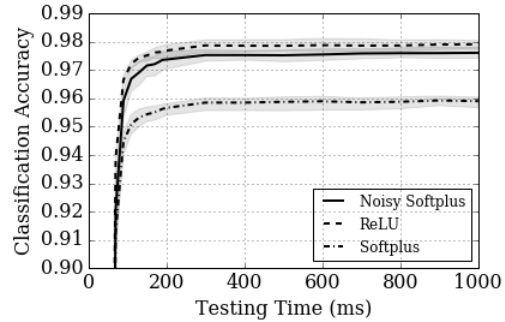
Table 1: Comparisons of classification accuracy (in %) of ANN-trained convolutional neural models on original DNN, NEST simulated SNN, and SNN with fine-tuned (FT) model.

| Trial No. | 1 | | | 2 | | | 3 | | |
|-----------------------|------|------|------|------|------|------|------|------|------|
| Model | DNN | SNN | FT | DNN | SNN | FT | DNN | SNN | FT |
| Noisy Softplus | 1.91 | 2.76 | 2.45 | 1.79 | 2.56 | 2.19 | 1.76 | 2.55 | 2.10 |
| ReLU | 1.36 | 2.03 | 1.88 | 1.46 | 2.28 | 2.00 | 1.36 | 2.25 | 2.12 |
| Softplus | 2.30 | 5.66 | 3.91 | 2.75 | 5.22 | 3.55 | 2.42 | 6.62 | 3.87 |

204 The best classification accuracy achieved by SNN was 98.85%, a 0.20% drop from ANN test (99.05%),
 205 which was trained with ReLU and fine-tuned by Noisy Softplus. The network structure was the
 206 same with the state-of-the-art model which reported the best classification accuracy of 99.1% [9] in
 207 ANN-trained SNNs: 12c5-2s-64c5-2s-10fc. Their nearly loss-less conversion from ANNs to SNNs
 208 was achieved by using IF neurons, while our network performs the best among SNNs consisted of
 209 LIF neurons to our knowledge.



(a) Before fine tuning



(b) After fine tuning.

Figure 8: The classification accuracy of 3 trials (averaged in bold lines, grey shading shows the range between minimum to maximum) over short response times, with (a) trained weights before fine tuning, and (b) after fine tuning.

210 As it is a major concern in neuromorphic vision, the recognition performance over short response
 211 times is also estimated in Figure 8. After fine tuning, Softplus significantly reduced the mismatch

212 since the randomness among the three trials shrinks to a range similar to other experiments. More
 213 obviously, fine tuning improved its classification accuracy and the response latency. Notice that all of
 214 the networks trained by three different activation functions showed a very similar stabilisation curve
 215 against time, which means they all reached an accuracy close to their best by only taking 300 ms of
 216 test.

217 3.3 Power Consumption

218 Noisy Softplus can easily be used for energy cost estimation for SNNs. For a single neuron, the
 219 energy consumption of the synaptic events it triggers is:

$$\begin{aligned} E_j &= \lambda_j N_j T E_{syn} \\ &= \frac{y_j N_j T E_{syn}}{\tau_{syn}}, \end{aligned} \quad (7)$$

220 where λ_j is the output firing rate, N_j is the number of post-synaptic neurons it connects to, T is
 221 the testing time, and E_{syn} is the energy cost for a synaptic event of some specific neuromorphic
 222 hardware, for example, about 8 nJ on SpiNNaker [17]. Thus to estimate the whole network, we can
 223 sum up all the synaptic events of all the neurons:

$$\sum_j E_j = \frac{T E_{syn}}{\tau_{syn}} \sum_j y_j N_j. \quad (8)$$

224 Thus, it may cost SpiNNaker 0.064 W, 192 J running for 3,000 s with synaptic events of $8 \times 10^6/s$
 225 to classify 10,000 images (300 ms each) with an accuracy of 98.02%. The best performance reported
 226 using the larger network may cost SpiNNaker 0.43 W operating synaptic event rate at $5.34 \times 10^7/s$,
 227 consume 4271.6 J to classify all the images for 1 s each.

228 4 Discussions

229 Most significantly, we proposed the Noisy Softplus activation function which accurately models
 230 response firing rate of LIF neurons and overcomes the drawbacks of Siegert units.

- 231 • Noisy Softplus takes account of time correlation of the noisy synaptic current, e.g. τ_{syn} ,
 232 which matches more closely to the actual response firing rate.
- 233 • Noisy Softplus can be applied easily to any training method, for example BP, thanks to its
 234 differentiability.
- 235 • the calculation on Noisy Softplus is no more than Softplus function, except for doubled
 236 computation on weighted sum of its input (net and σ in Equations 2 and 3), which is much
 237 more simplified than Siegert function.
- 238 • as one of the ReLU-liked activation function, the output firing rate seldom exceed the
 239 working range of a LIF neuron, for example the firing rates were around 0-200 Hz in the
 240 ConvNet model.
- 241 • the learning performance of Noisy Softplus is between Softplus and ReLU, which is sup-
 242 posed to outperform most of the other popular activation functions: for instance sigmoid.

243 Moreover, we proposed complete SNN modelling method by using artificial neurons of combined
 244 activation; this method can be generalised to activation units other than Noisy Softplus. The training
 245 of an SNN model is exactly the same as ANN training, and the trained weights can be directly used
 246 in SNN without any transformation. This method is simpler and even more straight-forward than the
 247 other ANN offline training methods which requires an extra step of converting ANN-trained weights
 248 to SNN's.

249 In terms of classification/recognition accuracy, the performance of ANN-trained SNNs is nearly
 250 equivalent as ANNs, and the performance loss can be partially solved by fine tuning. The best
 251 classification accuracy of 98.85% using LIF neurons in PyNN simulation outperforms state-of-the-art
 252 SNN models of LIF neurons which will be listed in Chapter 6, and is very close to the result using IF
 253 neurons [9].

Acknowledgments

To be added after reviewing.

References

- [1] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: a large-scale hierarchical image database. In *Computer Vision and Pattern Recognition, 2009. CVPR 2009. IEEE Conference on*, pages 248–255, 2009.
- [2] David Silver, Aja Huang, Chris J Maddison, Arthur Guez, Laurent Sifre, George Van Den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, et al. Mastering the game of go with deep neural networks and tree search. *Nature*, 529(7587):484–489, 2016.
- [3] Giacomo Indiveri, Elisabetta Chicca, and Rodney J Douglas. Artificial cognitive systems: from VLSI networks of spiking neurons to neuromorphic cognition. *Cognitive Computation*, 1(2):119–127, 2009.
- [4] F. Jug, J. Lengler, C. Krautz, and A. Steger. Spiking networks and their rate-based equivalents: does it make sense to use Siegert neurons? In *Swiss Soc. for Neuroscience*, 2012.
- [5] Evangelos Stamatias, Daniel Neil, Francesco Galluppi, Michael Pfeiffer, Shih-Chii Liu, and Steve Furber. Scalable energy-efficient, low-latency implementations of trained spiking deep belief networks on SpiNNaker. In *Neural Networks (IJCNN), 2015 International Joint Conference on*, pages 1–8. IEEE, 2015.
- [6] Steve B Furber, Francesco Galluppi, Sally Temple, Luis Plana, et al. The SpiNNaker project. *Proceedings of the IEEE*, 102(5):652–665, 2014.
- [7] Xavier Glorot, Antoine Bordes, and Yoshua Bengio. Deep sparse rectifier neural networks. In *International Conference on Artificial Intelligence and Statistics*, pages 315–323, 2011.
- [8] Yongqiang Cao, Yang Chen, and Deepak Khosla. Spiking deep convolutional neural networks for energy-efficient object recognition. *International Journal of Computer Vision*, 113(1):54–66, 2015.
- [9] Peter U Diehl, Daniel Neil, Jonathan Binas, Matthew Cook, Shih-Chii Liu, and Michael Pfeiffer. Fast-classifying, high-accuracy spiking deep networks through weight and threshold balancing. In *Neural Networks (IJCNN), 2015 International Joint Conference on*, pages 1–8. IEEE, 2015.
- [10] Nitish Srivastava, Geoffrey E Hinton, Alex Krizhevsky, Ilya Sutskever, and Ruslan Salakhutdinov. Dropout: a simple way to prevent neural networks from overfitting. *Journal of Machine Learning Research*, 15(1):1929–1958, 2014.
- [11] Peter U Diehl, Guido Zarrella, Andrew Cassidy, Bruno U Pedroni, and Emre Neftci. Conversion of artificial recurrent neural networks to spiking neural networks for low-power neuromorphic hardware. *arXiv preprint*, 2016.
- [12] Paul A Merolla, John V Arthur, Rodrigo Alvarez-Icaza, Andrew S Cassidy, Jun Sawada, Filipp Akopyan, Bryan L Jackson, Nabil Imam, Chen Guo, Yutaka Nakamura, et al. A million spiking-neuron integrated circuit with a scalable communication network and interface. *Science*, 345(6197):668–673, 2014.
- [13] Eric Hunsberger and Chris Eliasmith. Spiking deep networks with lif neurons. *arXiv preprint*, 2015.
- [14] Wulfram Gerstner and Werner Kistler. *Spiking Neuron Models: An Introduction*. Cambridge University Press, New York, NY, USA, 2002.
- [15] Qian Liu and Steve Furber. Noisy softplus: A biology inspired activation function. In *Proceedings of the 23rd International Conference on Neural Information Processing, Part IV*, pages 405–412, 2016.
- [16] Qian Liu, Evangelos Pineda-García Garibaldia nd Stamatias, Teresa Serrano-Gotarredona, and Steve Furber. Benchmarking spike-based visual recognition: A dataset and evaluation. *Frontiers in Neuroscience*, 10:496, 2016. <http://journal.frontiersin.org/article/10.3389/fnins.2016.00496>.
- [17] Evangelos Stamatias, Francesco Galluppi, Cameron Patterson, and Steve Furber. Power analysis of large-scale, real-time neural networks on SpiNNaker. In *Neural Networks (IJCNN), The 2013 International Joint Conference on*, pages 1–8, 2013.